

## **Ameya AI – Hands-On Workshop Facilitator Master Guide**

This document is the single source of truth for facilitation. It is designed for reference during the session and can be shared with participants AFTER completion.

Principles:

- No slides
- No theory
- Hands-on only
- Local-first, enterprise-shaped AI

### **Open-Source Models for Local AI (New Section)**

This section explains where to find open-source models, how to choose the right model for a local machine, and how Ollama Local differs from Ollama Cloud. This content is critical context for Session 2.

#### **Where to Find Open-Source Models**

Ollama:

- Simplest way to run modern LLMs locally
- Handles model lifecycle, quantization, and serving
- Exposes a local HTTP API (localhost:11434)
- Best choice for workshops and enterprise PoCs

Hugging Face:

- Largest open-source model hub
- Useful for advanced customization and fine-tuning
- Requires more setup and ML expertise

#### **How to Choose a Model for Local Execution**

Choose models based on constraints, not benchmarks:

- Hardware: CPU-only vs GPU
- Latency: interactive vs batch
- Context length: resumes vs large documents
- Accuracy vs speed trade-offs
- Compliance: data must stay local

Recommended defaults:

- llama3.2 for CPU-only machines

- mistral / qwen variants if more compute is available

### **Model Size vs Laptop Reality**

Practical guidance:

- 7B models → safe for most laptops (CPU)
- 13B models → slower, usable with patience
- 30B+ → not suitable for this workshop

Workshop rule:

- If your laptop struggles, downgrade the model.
- Reliability > raw intelligence.

### **Ollama Local vs Ollama Cloud**

Ollama Local:

- Fully local inference
- Zero data egress
- Predictable cost
- Enterprise and regulated workloads

Ollama Cloud:

- Managed inference
- Higher throughput
- Data leaves local boundary
- Best for non-sensitive demos

Workshop Rule: Ollama Local ONLY.

### **Running Models Locally**

ollama pull llama3.2

ollama run llama3.2

## **1. Workshop Overview**

Session 1: AI-Native Development with Agentic IDEs (60 mins)

Session 2: Local LLMs and Enterprise AI (60 mins)

IDE: Cursor ONLY

LLM: Local LLM via Ollama

## 2. One-Page Facilitator Checklist

- Cursor IDE ready
- Python 3.10+
- Ollama running
- Repo cloned and tested

## 3. README – Exact Commands

```
git clone <repo-url>
cd ai-native-skill-platform
python -m venv venv
source venv/bin/activate
pip install -r requirements.txt
```

```
uvicorn app.main:app --reload
```

## 4. Facilitator-Only Prompt Sheet

### Prompt 1 – Service Scaffolding

*You are an AI engineer.*

*Create a FastAPI service that uses a LOCAL LLM via Ollama (<http://localhost:11434>).*

*Do NOT use any cloud LLM APIs or API keys.*

*Endpoints:*

- *GET /health*
- *POST /upload\_resume*
- *POST /analyze\_skills*

*Only code. No explanations.*

### Prompt 2 – Local LLM + Contract Enforcement

*Extend the existing FastAPI service.*

*Constraints:*

- *Use Python requests or httpx to call Ollama (/api/generate or /api/chat)*
- *Model name must come from env var OLLAMA\_MODEL (default: llama3.2)*
- */upload\_resume must accept text or file and store it under ./data/resumes*

- */analyze\_skills must call Ollama and return structured JSON*

*Return JSON format:*

```
{  
  "skills": [],  
  "years_experience": "optional",  
  "roleSuggestions": []  
}
```

*Add basic error handling if Ollama is not running.*

*Only code. No explanations.*

### Prompt 3 – Skill AI Intelligence Layer

*Extend the service with intelligence features.*

*Add:*

- *skill\_gap\_analysis(target\_role)*
- *weekly\_learning\_task\_generator()*

*Behavior:*

- *Compare extracted skills against the target role*
- *Identify missing skills*
- *Generate 5–7 practical weekly learning tasks per missing skill*

*Return a single structured JSON response combining:*

- *extracted\_skills*
- *missing\_skills*
- *weekly\_tasks*

*Fix any issues automatically.*

*Only code. No explanations.*

## 5. Further improvements and assignments to participants:

- Modularize the codebase (agents, llm, api layers)
- Add support for PDF and Word resumes
- Add simple logging for each request
- Add a basic token or rate limit
- Replace in-memory storage with a vector database

## 6.What Participants Should Walk Away With

By the end of this workshop, participants will have:

- Built an AI-native service using an agentic IDE
- Used a local LLM with enterprise boundaries
- Understood how skills, gaps, and tasks can be modeled as systems
- Seen how agents + local models form the foundation of real AI platforms

## 7. Day-2 Enterprise Extension Roadmap

1. Auth & RBAC
2. Observability (logs, metrics)
3. Persistent vector storage
4. Docker & secure deployment
5. Multi-agent orchestration